Traffic Sign Recognition using Deep Learning

Zaid Khan Final Project Presentation

Overview of Traffic Sign Recognition

- Recognizing traffic signs is crucial for autonomous driving.
- Real-time detection and classification enhance safety for automated and driver-assist systems.

https://www.tesla.com/autopilot

Why?

- Relevance to real-world applications in autonomous vehicles
  - *tesla please stop rejecting my internship apps.*
- Opportunity to learn more about CNN architecture and live implementation.
- Challenge of working with a dataset with multiple classes and varied conditions.
- The German Traffic Sign Recognition Benchmark (GTSRB) already exists as a dataset on kaggle.

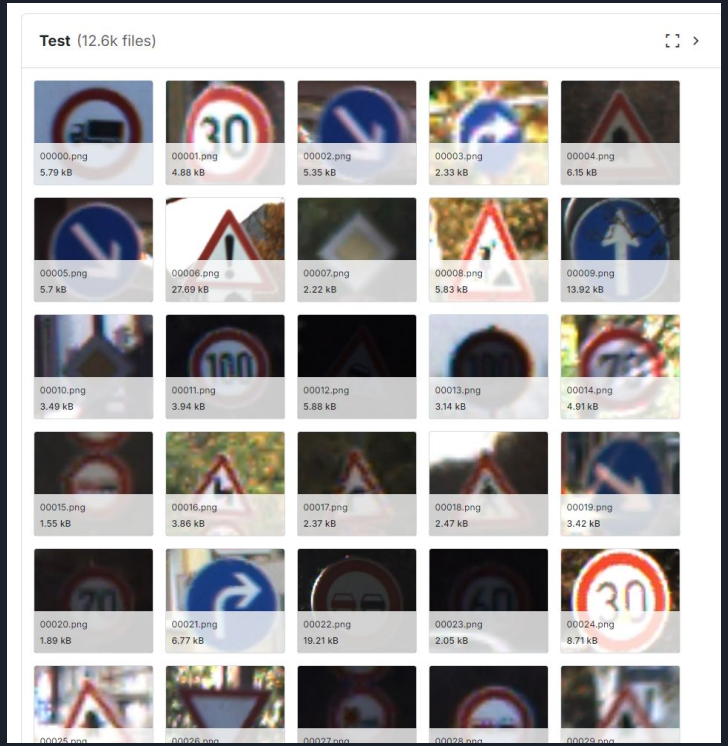

[GTSRB Dataset on Kaggle]

The Model

- Libraries and Dataset:
  - Torch, Torchvision
  - Dataset: GTSRB, consisting of 43 traffic sign classes.
- Model Architecture:
  - Using ResNet50 as my base case for the model.
  - Using ResNet-152 to experiment between the two.
  - Adjusting the final layer for 43-class classification.

Training

- Training Loop:
  - Model parameters are optimized using Adam and cross-entropy loss.
  - Accuracy and loss computed for both training and validation sets each epoch.

- Hyperparameters:
  - Batch size: 32, Epochs: 50, Learning Rate: 0.001.
- Used RTX 4060ti for cuda processing

```
Epoch: 40, Loss: 0.01579, Accuracy: 93.32542%
Epoch: 41, Loss: 0.01554, Accuracy: 91.71021%
Epoch: 42, Loss: 0.01627, Accuracy: 91.85273%
Epoch: 43, Loss: 0.01664, Accuracy: 89.39826%
Epoch: 44, Loss: 0.01241, Accuracy: 92.41489%
Epoch: 45, Loss: 0.02167, Accuracy: 89.31116%
Epoch: 46, Loss: 0.01460, Accuracy: 93.27791%
Epoch: 47, Loss: 0.01723, Accuracy: 92.66825%
Epoch: 48, Loss: 0.01810, Accuracy: 91.05305%
Epoch: 49, Loss: 0.01952, Accuracy: 90.49089%
Epoch: 50, Loss: 0.01759, Accuracy: 92.64450%
Total Training Time: 1104.11 seconds
Testing Time: 18.40 minutes
Accuracy of the network on the test images: 92.64%
Testing Time: 4.58 seconds
Total execution time: 1113.31 seconds
Testing Time: 18.56 minutes
```

Real Time Recognition

- Video Capture and Preprocessing:
  - Capturing frames from a video file, resizing, and transforming each frame for model input.

- Inference:
  - Softmax to obtain probability for each class, with thresholding to determine detection confidence.
  - Displaying results on video frames in real-time.

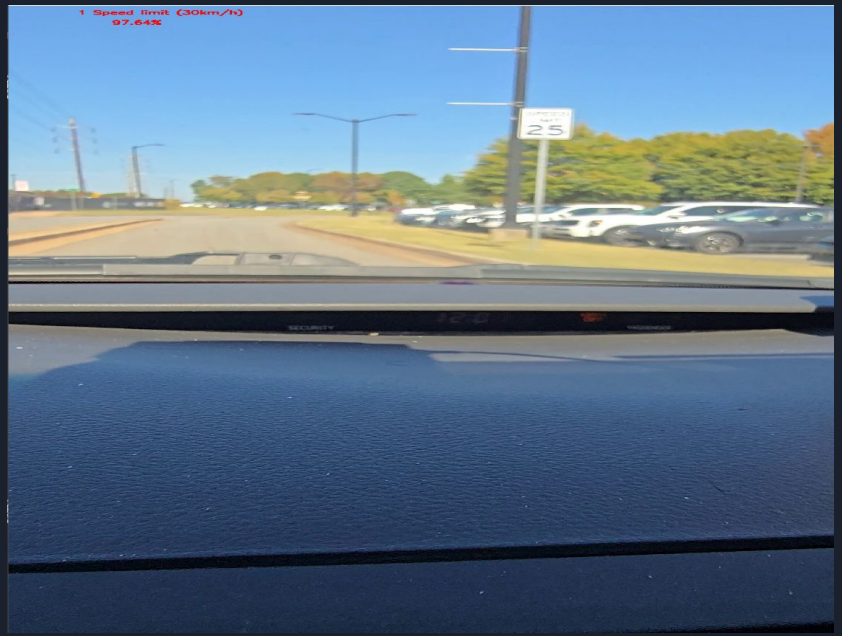

Challenges

- Noise
  - The video I had was very blurry because my camera kept refocusing on my dashboard

- General model accuracy on the video feed
  - The model was trained on a foreign dataset
  - Different batch sizes, models, and general hyperparameters could be affecting the classification
  - The recognition model looks at the entire image not at the specific signs

Future work

- Better video for testing
- Create a detection algorithm to detect where the traffic signs are in the video
- More localized datasets
  - https://www.kaggle.com/datasets/pkdarabi/cardetection
  - https://git-disl.github.io/GTDLBench/datasets/lisa_traffic_sign_dataset/
- Finally move on to flags and racetrack machine vision classifications